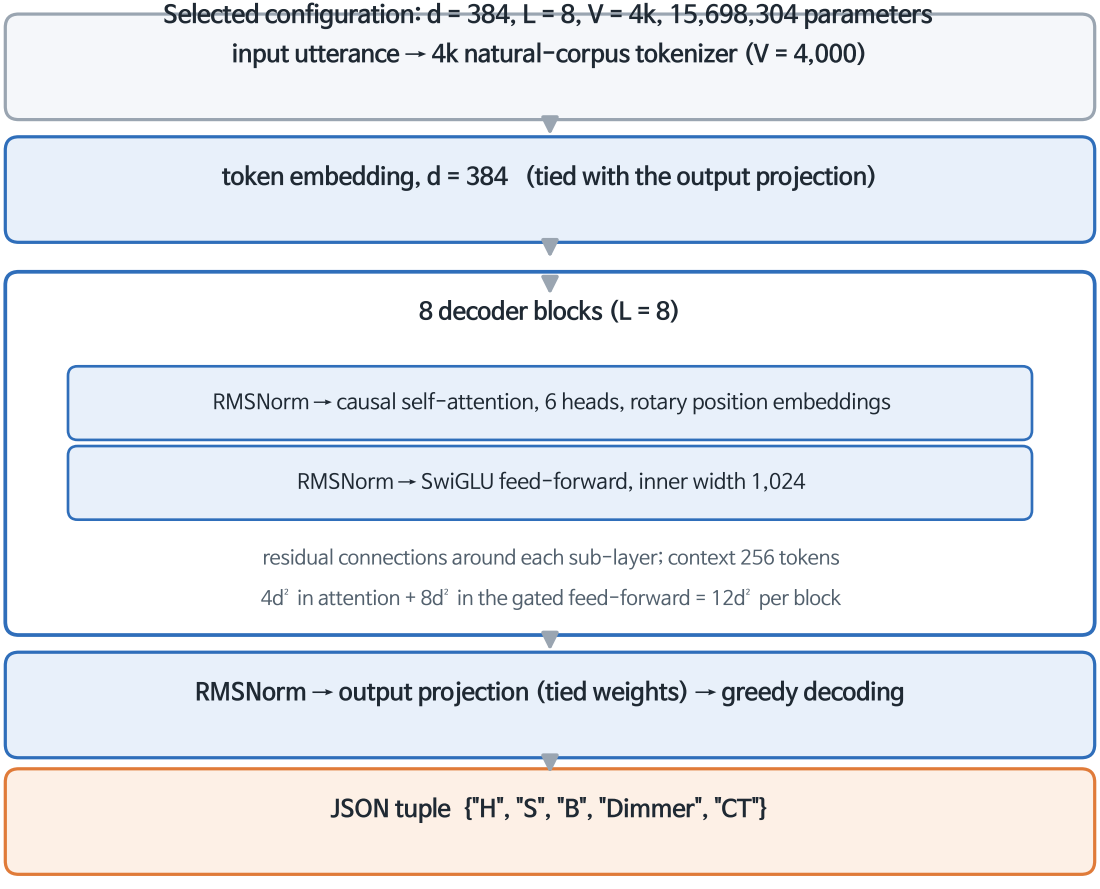
Selected configuration: d = 384, L = 8, V = 4k, 15,698,304 parameters
input utterance → 4k natural-corpus tokenizer (V = 4,000)
token embedding, d = 384 (tied with the output projection)
8 decoder blocks (L = 8)
RMSNorm → causal self-attention, 6 heads, rotary position embeddings
RMSNorm → SwiGLU feed-forward, inner width 1,024
residual connections around each sub-layer; context 256 tokens
4d² in attention + 8d² in the gated feed-forward = 12d² per block
RMSNorm → output projection (tied weights) → greedy decoding
JSON tuple {"H", "S", "B", "Dimmer", "CT"}